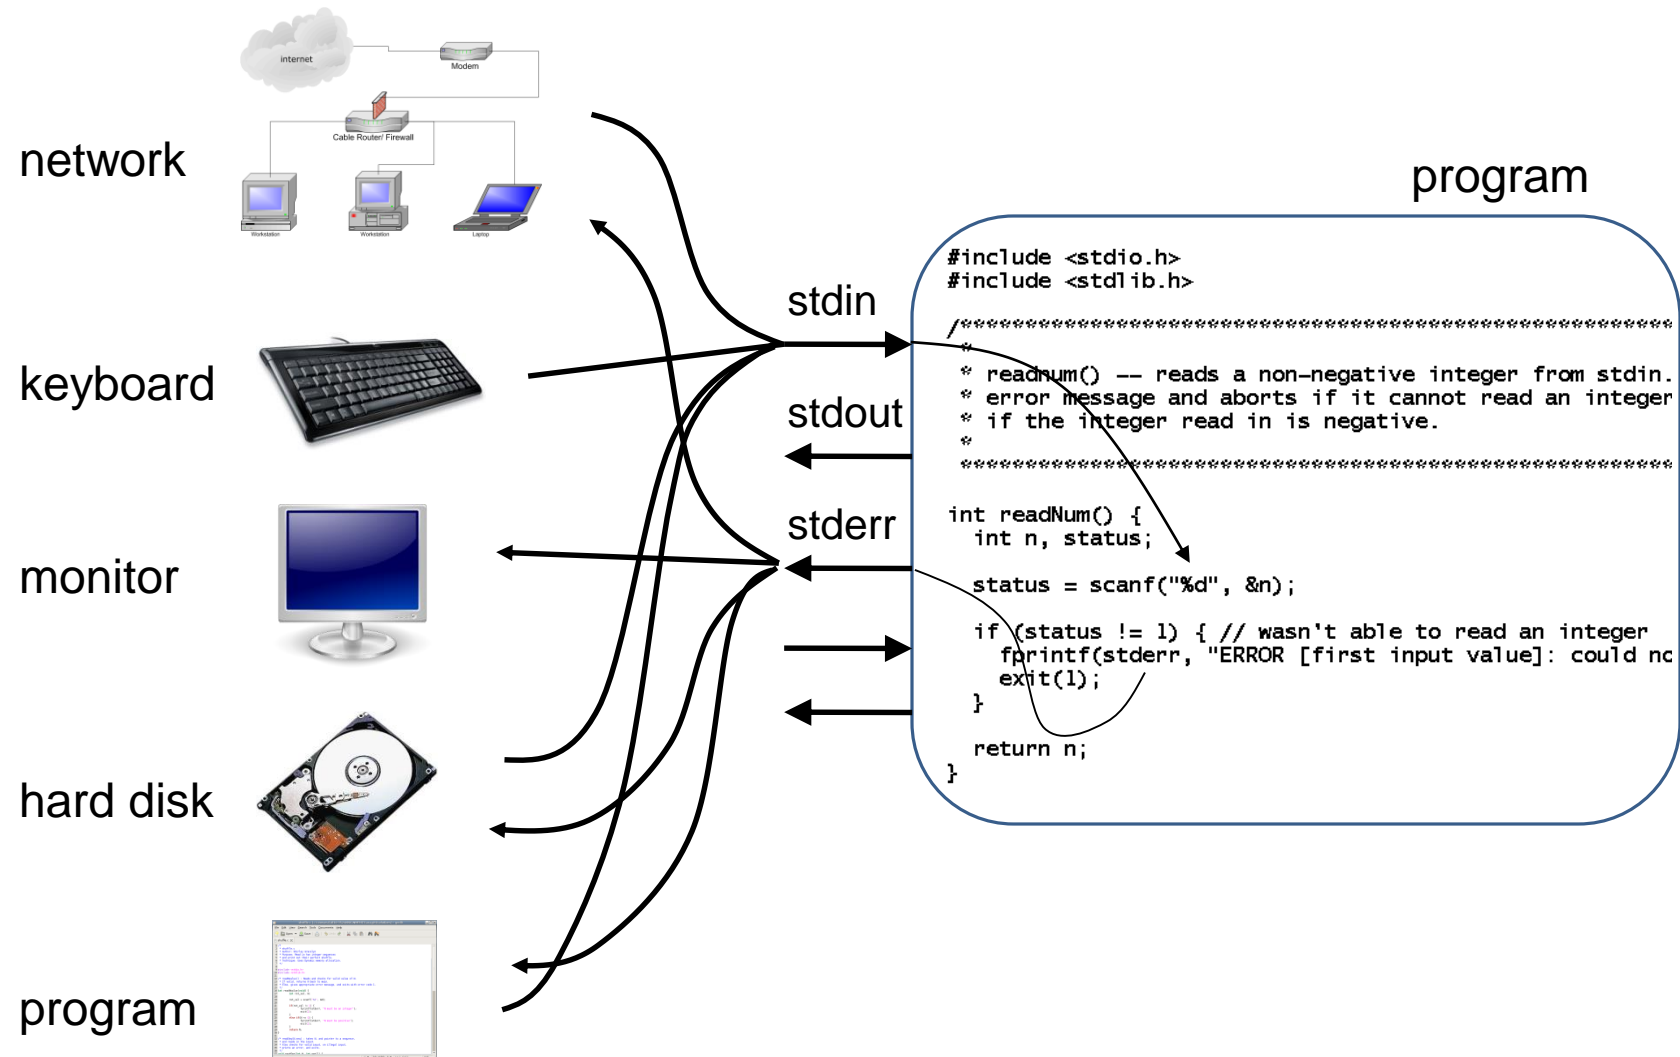


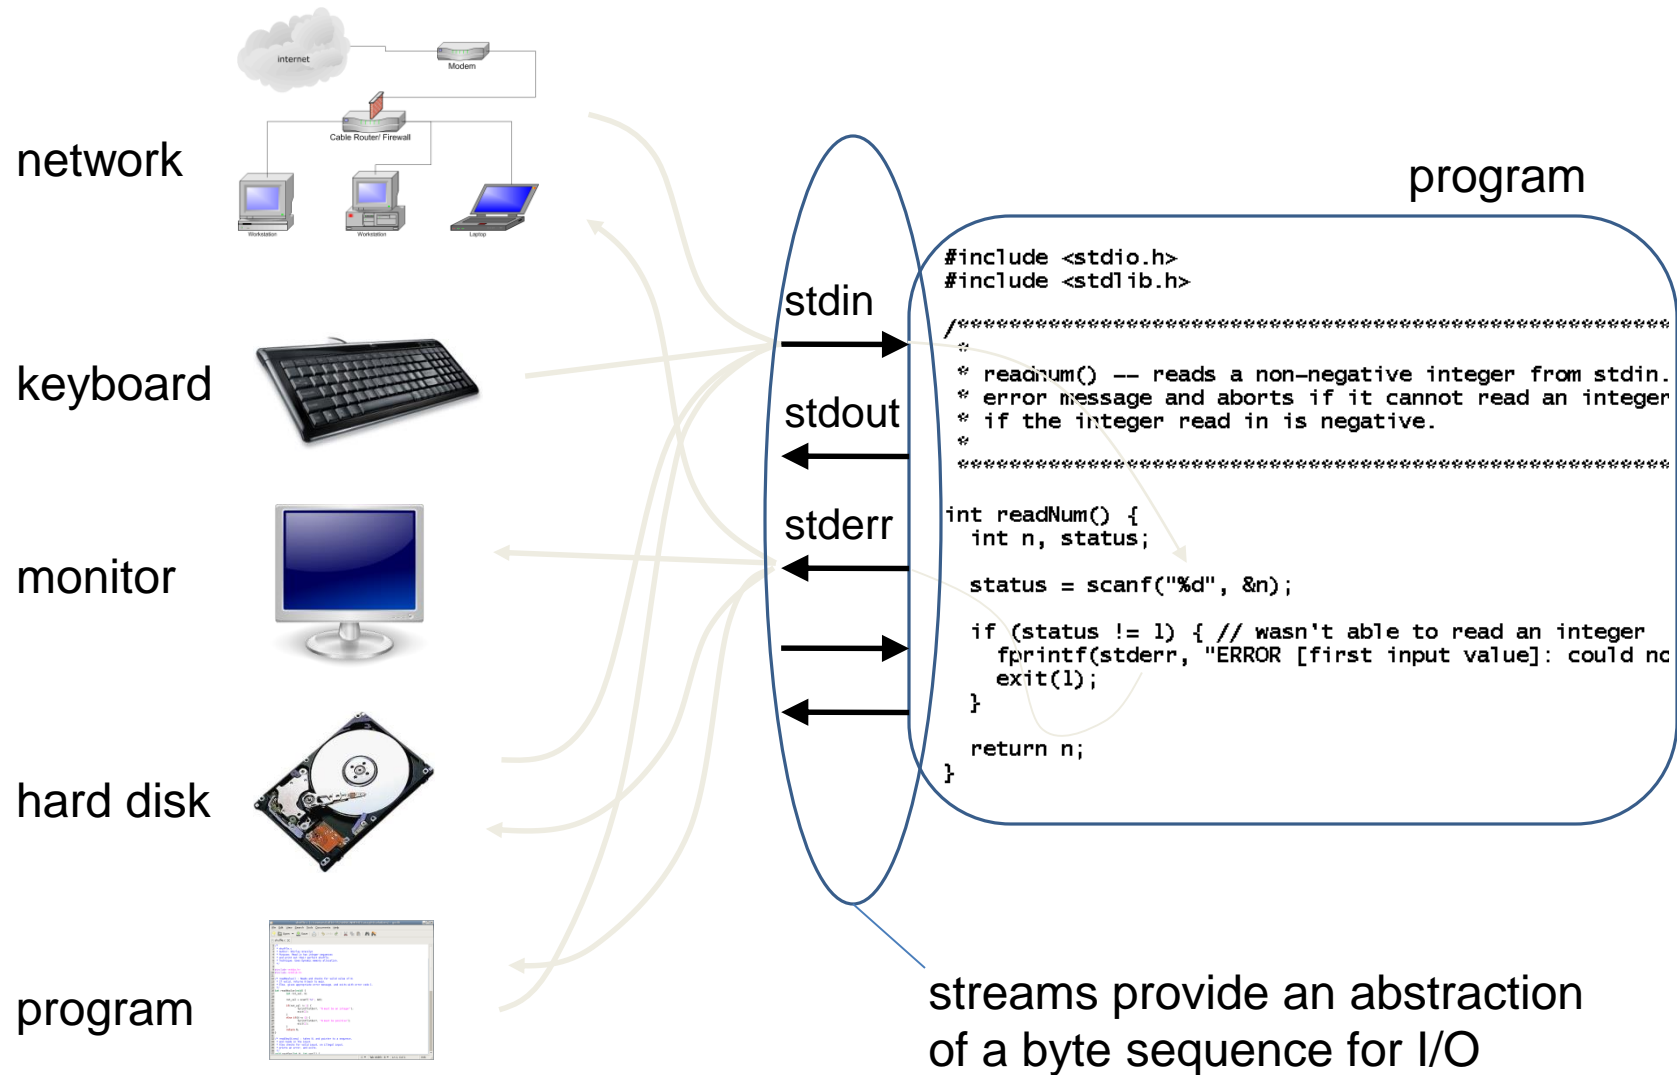
Streams

- A *stream* is any source of input or any destination for output
 - conceptually, just a sequence of bytes
 - accessed through a file pointer, which has type
FILE *
 - however, not all streams are associated with files
 - three standard predefined streams: stdin, stdout, stderr

Streams



Streams



Typical structure of I/O operations

A program's I/O operations usually have the following structure:

1. Open a file

`fopen`

2. Perform I/O

`fprintf, fscanf`
`fread, fwrite`
`fgets`

3. Close the file

`fclose`

Opening a file

FILE * fopen(char *filename, char *mode)

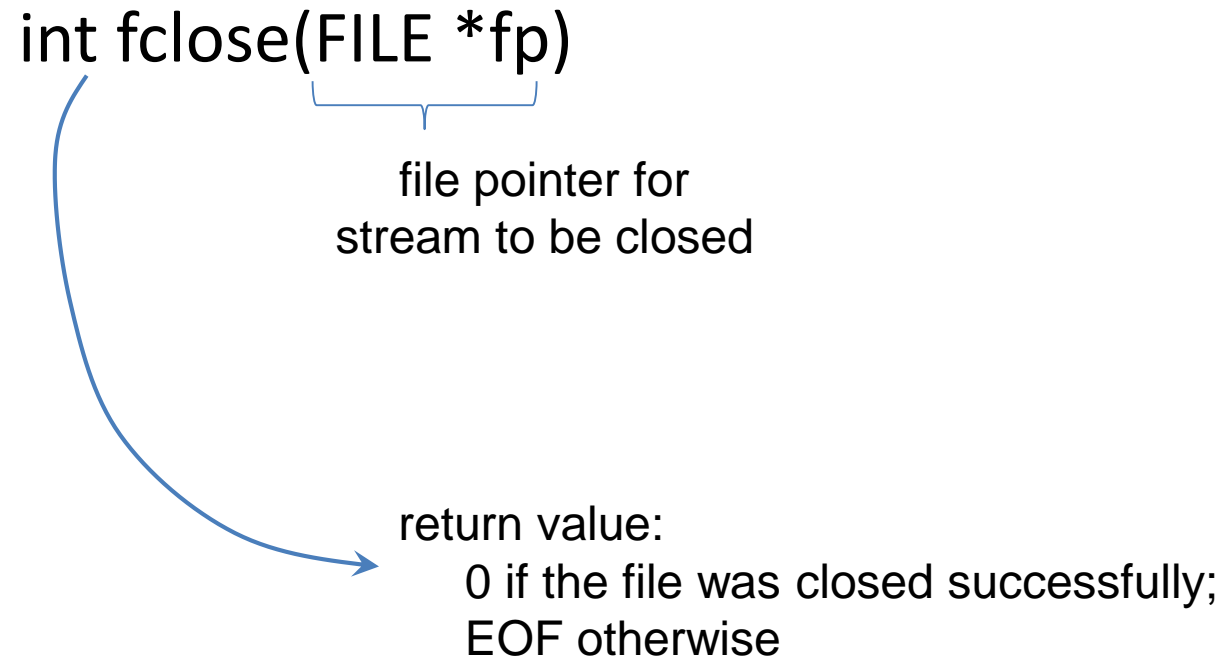
name of file to open

file pointer for the stream, if
fopen succeeds;
NULL otherwise

"r"	read
"w"	write (file need not exist)
"a"	append (file need not exist)
"r+"	read and write, starting at the beginning
"w+"	read and write; truncate file if it exists
"a+"	read and write; append if file exists

Closing a file

`int fclose(FILE *fp)`



file pointer for
stream to be closed

return value:
0 if the file was closed successfully;
EOF otherwise

Code structure

```
FILE *fp;
```

```
...
```

```
fp = fopen(filename, "r");  // or whatever mode is appropriate
```

```
if (fp == NULL) {
```

```
    ... give error message and exit ...
```

```
}
```

```
... read and process file ...
```

```
status = fclose(fp);
```

```
if (status == EOF) {
```

```
    ... give error message...
```

```
}
```

Reading and writing

- **fprintf, fscanf**
 - similar to printf and scanf, with additional **FILE *** argument
- **fread(*ptr*, *sz*, *num*, *fp*)**
 - reads *num* elements, each of size *sz*, from stream *fp* and stores them at *ptr*
 - does not distinguish between end-of-file and error
 - use **feof()** and **ferror()**
- **fwrite(*ptr*, *sz*, *num*, *fp*)**
 - writes *num* elements, of size *sz*, from *ptr* into stream *fp*
- return values:
 - no. of items successfully read/written (not no. of bytes)

Example: fread

The image shows a code editor window titled "hed: /cs/www/classes/cs352/spring10/Code/ex.6.FileIO" containing the source code for a C program named `nlines.c`. The code is annotated with four callout boxes connected by red lines, explaining key steps in the program's execution:

- open the file for reading**: Points to the line `fp = fopen(argv[1], "r");`.
- check for EOF (fread does not distinguish between EOF and error)**: Points to the condition `while (!feof(fp))` in the loop header.
- read from the file (as much as will fit in the buffer)**: Points to the `fread` call: `nchars = fread(buf, sizeof(char), 4096/sizeof(char), fp);`.
- close the file when done (should check for errors)**: Points to the line `fclose(fp);`.

```
% cat nlines.c
/*
 * File: nlines.c
 * Purpose: count the no. of lines in a given file; illustrate file I/O
 */
#include <stdio.h>

char buf[4096];
int nlines = 0;

int main(int argc, char *argv[]) {
    FILE *fp;
    int i, nchars;

    if (argc < 2) {
        fprintf(stderr, "Usage: %s file\n", argv[0]);
        return 1;
    }

    fp = fopen(argv[1], "r");
    if (fp == NULL) {
        fprintf(stderr, "Could not open file %s\n", argv[1]);
        return 1;
    }

    while (!feof(fp)) {
        nchars = fread(buf, sizeof(char), 4096/sizeof(char), fp);

        for (i = 0; i < nchars; i++) {
            if (buf[i] == '\n') {
                nlines += 1;
            }
        }
    }

    printf("File %s contains %d lines\n", argv[1], nlines);
    fclose(fp);
    return 0;
}
```

Example: fscanf

```
hed: /cs/www/classes/cs352/spring10/Code/ex.6.FileIO
% cat add-numbers-in-file.c
/*
 * File: add-numbers-in-file.c
 * Purpose: add up the integers in a given file; illustrate file I/O
 */
#include <stdio.h>

int main(int argc, char *argv[]) {
    FILE *fp;
    int status, num, sum = 0;

    if (argc < 2) {
        fprintf(stderr, "Usage: %s file\n", argv[0]);
        return 1;
    }

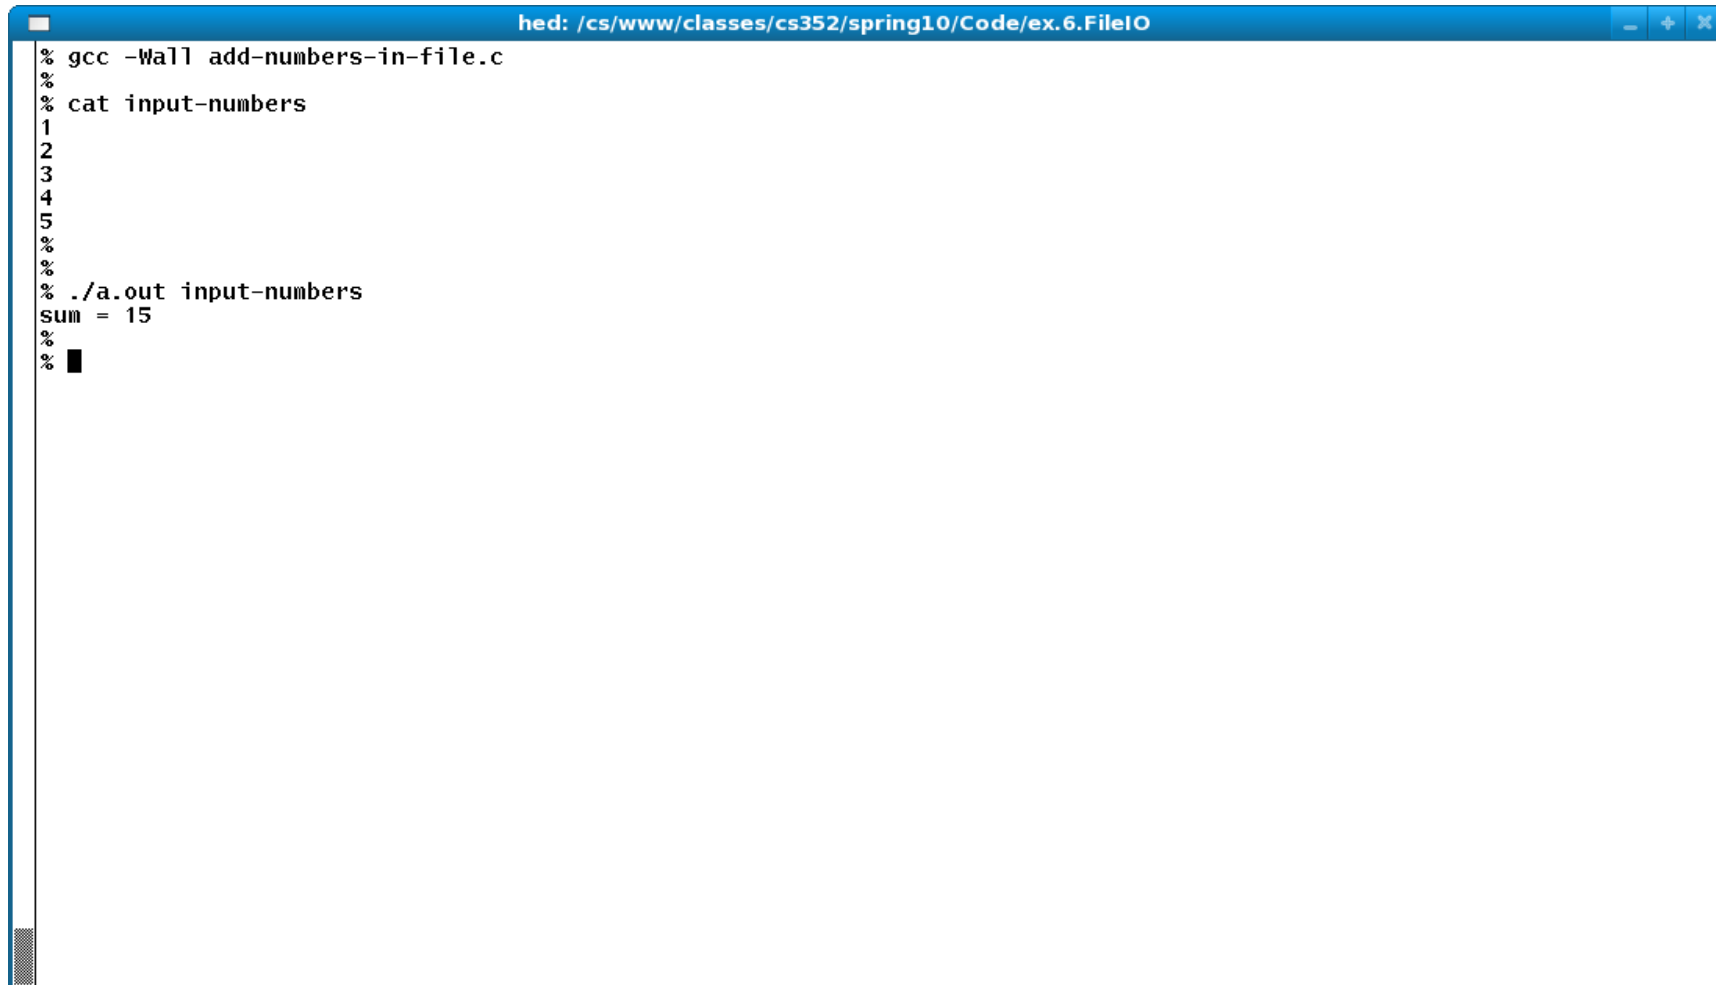
    fp = fopen(argv[1], "r");
    if (fp == NULL) {
        fprintf(stderr, "Could not open file %s\n", argv[1]);
        return 1;
    }

    while ((status = fscanf(fp, "%d", &num)) != EOF) {
        if (status == 0) {
            fprintf(stderr, "Non-numeric data in input file %s\n", argv[1]);
            return 1;
        }

        sum += num;
    }

    printf("sum = %d\n", sum);
    fclose(fp);
    return 0;
}
% █
```

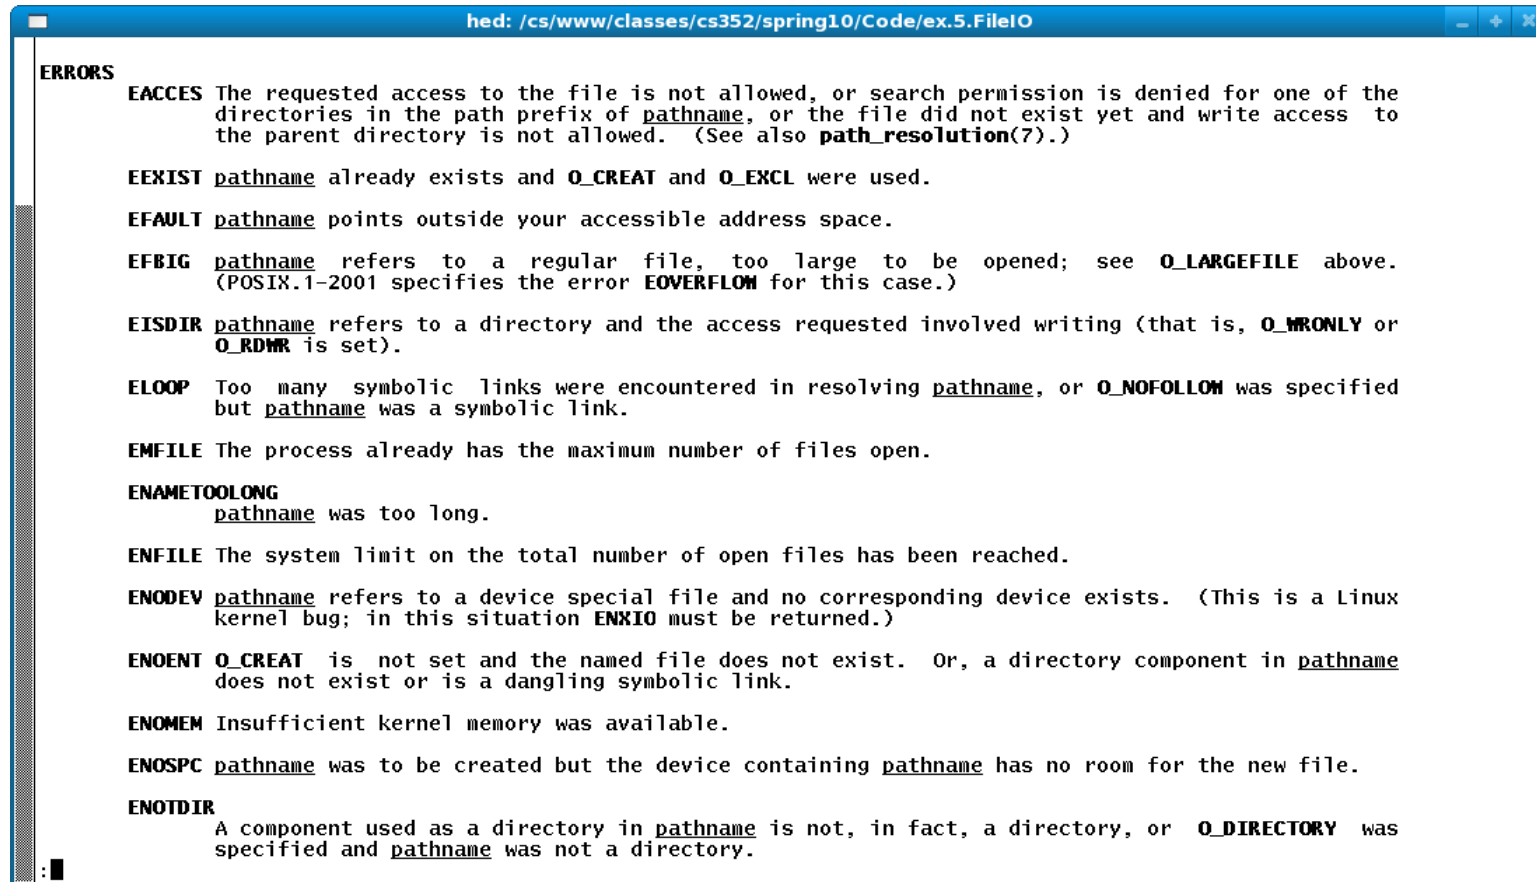
fscanf



```
hed: /cs/www/classes/cs352/spring10/Code/ex.6.FileIO
% gcc -Wall add-numbers-in-file.c
%
% cat input-numbers
1
2
3
4
5
%
%
% ./a.out input-numbers
sum = 15
%
% █
```

Errors

- `fopen()` can fail for many different reasons



```
hed: /cs/www/classes/cs352/spring10/Code/ex.5.FileIO
ERRORS
EACCES The requested access to the file is not allowed, or search permission is denied for one of the
        directories in the path prefix of pathname, or the file did not exist yet and write access to
        the parent directory is not allowed. (See also path_resolution(7).)

EEXIST pathname already exists and O_CREAT and O_EXCL were used.

EFAULT pathname points outside your accessible address space.

EFBIG pathname refers to a regular file, too large to be opened; see O_LARGEFILE above.
        (POSIX.1-2001 specifies the error E_OVERFLOW for this case.)

EISDIR pathname refers to a directory and the access requested involved writing (that is, O_WRONLY or
        O_RDWR is set).

ELOOP Too many symbolic links were encountered in resolving pathname, or O_NOFOLLOW was specified
        but pathname was a symbolic link.

EMFILE The process already has the maximum number of files open.

ENAMETOOLONG
        pathname was too long.

ENFILE The system limit on the total number of open files has been reached.

ENODEV pathname refers to a device special file and no corresponding device exists. (This is a Linux
        kernel bug; in this situation ENXIO must be returned.)

ENOENT O_CREAT is not set and the named file does not exist. Or, a directory component in pathname
        does not exist or is a dangling symbolic link.

ENOMEM Insufficient kernel memory was available.

ENOSPC pathname was to be created but the device containing pathname has no room for the new file.

ENOTDIR
        A component used as a directory in pathname is not, in fact, a directory, or O_DIRECTORY was
        specified and pathname was not a directory.
```

Giving sensible error messages: perror

```
hed: /cs/www/classes/cs352/spring10/Code/ex.5.FileIO
PERROR(3)          Linux Programmer's Manual          PERROR(3)

NAME
    perror - print a system error message

SYNOPSIS
    #include <stdio.h>

    void perror(const char *s);

    #include <errno.h>

    const char *sys_errlist[];
    int sys_nerr;
    int errno;

Feature Test Macro Requirements for glibc (see feature_test_macros(7)):

    sys_errlist, sys_nerr: _BSD_SOURCE

DESCRIPTION
    The routine perror() produces a message on the standard error output, describing the last error encountered during a call to a system or library function. First (if s is not NULL and *s is not a null byte ('\0')) the argument string s is printed, followed by a colon and a blank. Then the message and a new-line.

    To be of most use, the argument string should include the name of the function that incurred the error. The error number is taken from the external variable errno, which is set when errors occur but not cleared when non-erroneous calls are made.

    The global error list sys_errlist[] indexed by errno can be used to obtain the error message without the newline. The largest message number provided in the table is sys_nerr - 1. Be careful when directly accessing this list because new error values may not have been added to sys_errlist[].

    When a system call fails, it usually returns -1 and sets the variable errno to a value describing what went wrong. (These values can be found in <errno.h>.) Many library functions do likewise. The function perror() serves to translate this error code into human-readable form. Note that errno is undefined after a successful library call: this call may well change this variable, even though it succeeds, for example because it internally used some other library function that failed. Thus, if a
```

perror

```
hed: /cs/www/classes/cs352/spring10/Code/ex.6.FileIO
* File: nlines-perror.c
* Purpose: count the no. of lines in a given file; illustrate file I/O
* as well as the use of perror() to give sensible error messages.
*/
#include <stdio.h>
#include <errno.h>

char buf[4096];
int nlines = 0;
int errno;

int main(int argc, char *argv[]) {
    FILE *fp;
    int i, nchars;

    if (argc < 2) {
        fprintf(stderr, "Usage: %s file\n", argv[0]);
        return 1;
    }

    fp = fopen(argv[1], "r");
    if (fp == NULL) {
        perror(argv[1]); //WAS: fprintf(stderr, "Could not open file %s\n", argv[1]);
        return 1;
    }

    while (!feof(fp)) {
        nchars = fread(buf, sizeof(char), 4096/sizeof(char), fp);
        for (i = 0; i < nchars; i++) {
            if (buf[i] == '\n') {
                nlines += 1;
            }
        }
    }

    printf("File %s contains %d lines\n", argv[1], nlines);
    fclose(fp);
    return 0;
}
```

perror

```
hed: /cs/www/classes/cs352/spring10/Code/ex.6.FileIO
% gcc -Wall nlines-perror.c
%
% ls -al
total 28
drwxr-xr-x  2 debray dept 4096 2010-02-22 09:52 ./
drwxr-xr-x 10 debray dept 4096 2010-02-19 16:04 ../
-rwxr-xr-x  1 debray dept 7879 2010-02-22 09:52 a.out*
--w-----  1 debray dept   0 2010-02-22 09:45 FILE-WITH-NO-READ-PERMISSIONS
-rw-r--r--  1 debray dept  35 2010-02-19 16:47 infile0
-rw-r--r--  1 debray dept 704 2010-02-19 16:50 nlines.c
-rw-r--r--  1 debray dept 830 2010-02-22 09:47 nlines-perror.c
%
% ./a.out ./NONEXISTENT-FILE
./NONEXISTENT-FILE: No such file or directory
%
% ./a.out FILE-WITH-NO-READ-PERMISSIONS
FILE-WITH-NO-READ-PERMISSIONS: Permission denied
%
% █
```